

Planning Stage Report – Group 4 | STAT 301 (2025SS)

Student: Krishaant Pathmanathan

Dataset: Customer Personality Analysis

TA: Yian Lin

```
In [125... # install.packages("skimr")
# install.packages("ggcorrplot")
library(tidyverse)
library(skimr)
library(dplyr)
library(tidyr)
library(tibble)
library(ggplot2)
library(GGally)
library(dplyr)
library(janitor)
library(patchwork)
library(cowplot)
library(corrplot)
library(ggcorrplot)
```

```
In [126... df <- read_delim("marketing_campaign.csv", delim = "\t")
glimpse(df) # to see what the data looks like
skim(df) # to get a summary of the data
```

Rows: 2240 **Columns:** 29

— Column specification —

Delimiter: "\t"

chr (3): Education, Marital_Status, Dt_Customer

dbl (26): ID, Year_Birth, Income, Kidhome, Teenhome, Recency, MntWines, MntF...

i Use `spec()` to retrieve the full column specification for this data.

i Specify the column types or set `show_col_types = FALSE` to quiet this message.

Rows: 2,240

Columns: 29

\$ ID	<dbl> 5524, 2174, 4141, 6182, 5324, 7446, 965, 6177, 485...
\$ Year_Birth	<dbl> 1957, 1954, 1965, 1984, 1981, 1967, 1971, 1985, 19...
\$ Education	<chr> "Graduation", "Graduation", "Graduation", "Graduat...
\$ Marital_Status	<chr> "Single", "Single", "Together", "Together", "Married"
\$ Income	<dbl> 58138, 46344, 71613, 26646, 58293, 62513, 55635, 3...
\$ Kidhome	<dbl> 0, 1, 0, 1, 1, 0, 0, 1, 1, 1, 1, 0, 0, 1, 0, 0, 1, ...
\$ Teenhome	<dbl> 0, 1, 0, 0, 0, 1, 1, 0, 0, 1, 0, 0, 0, 1, 0, 0, 1, ...
\$ Dt_Customer	<chr> "04-09-2012", "08-03-2014", "21-08-2013", "10-02-2012", ...
\$ Recency	<dbl> 58, 38, 26, 26, 94, 16, 34, 32, 19, 68, 11, 59, 82...
\$ MntWines	<dbl> 635, 11, 426, 11, 173, 520, 235, 76, 14, 28, 5, 6, ...
\$ MntFruits	<dbl> 88, 1, 49, 4, 43, 42, 65, 10, 0, 0, 5, 16, 61, 2, ...
\$ MntMeatProducts	<dbl> 546, 6, 127, 20, 118, 98, 164, 56, 24, 6, 6, 11, 4, ...
\$ MntFishProducts	<dbl> 172, 2, 111, 10, 46, 0, 50, 3, 3, 1, 0, 11, 225, 3...
\$ MntSweetProducts	<dbl> 88, 1, 21, 3, 27, 42, 49, 1, 3, 1, 2, 1, 112, 5, 1...
\$ MntGoldProds	<dbl> 88, 6, 42, 5, 15, 14, 27, 23, 2, 13, 1, 16, 30, 14...
\$ NumDealsPurchases	<dbl> 3, 2, 1, 2, 5, 2, 4, 2, 1, 1, 1, 1, 1, 3, 1, 1, 3, ...
\$ NumWebPurchases	<dbl> 8, 1, 8, 2, 5, 6, 7, 4, 3, 1, 1, 2, 3, 6, 1, 7, 3, ...
\$ NumCatalogPurchases	<dbl> 10, 1, 2, 0, 3, 4, 3, 0, 0, 0, 0, 0, 4, 1, 0, 6, 0...
\$ NumStorePurchases	<dbl> 4, 2, 10, 4, 6, 10, 7, 4, 2, 0, 2, 3, 8, 5, 3, 12, ...
\$ NumWebVisitsMonth	<dbl> 7, 5, 4, 6, 5, 6, 6, 8, 9, 20, 7, 8, 2, 6, 8, 3, 8...
\$ AcceptedCmp3	<dbl> 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, ...
\$ AcceptedCmp4	<dbl> 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, ...
\$ AcceptedCmp5	<dbl> 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, ...
\$ AcceptedCmp1	<dbl> 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, ...
\$ AcceptedCmp2	<dbl> 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, ...
\$ Complain	<dbl> 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, ...
\$ Z_CostContact	<dbl> 3, 3, 3, 3, 3, 3, 3, 3, 3, 3, 3, 3, 3, 3, 3, 3, 3, ...
\$ Z_Revenue	<dbl> 11, 11, 11, 11, 11, 11, 11, 11, 11, 11, 11, 11, 11, 11, 11, 11, ...
\$ Response	<dbl> 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, ...

— Data Summary —

	Values
Name	df
Number of rows	2240
Number of columns	29

Column type frequency:	
character	3
numeric	26

Group variables	None
-----------------	------
















— Variable type: character —

	skim_variable	n_missing	complete_rate	min	max	empty	n_unique	whitespace
1	Education	0	1	3	10	0	5	0
2	Marital_Status	0	1	4	8	0	8	0
3	Dt_Customer	0	1	10	10	0	663	0

— Variable type: numeric —

	skim_variable	n_missing	complete_rate	mean	sd	p0
1	ID	0	1	5592.	3247.	0
2	Year_Birth	0	1	1969.	12.0	1893
3	Income	24	0.989	52247.	25173.	1730
4	Kidhome	0	1	0.444	0.538	0
5	Teenhome	0	1	0.506	0.545	0
6	Recency	0	1	49.1	29.0	0
7	MntWines	0	1	304.	337.	0
8	MntFruits	0	1	26.3	39.8	0
9	MntMeatProducts	0	1	167.	226.	0
10	MntFishProducts	0	1	37.5	54.6	0
11	MntSweetProducts	0	1	27.1	41.3	0
12	MntGoldProds	0	1	44.0	52.2	0
13	NumDealsPurchases	0	1	2.33	1.93	0
14	NumWebPurchases	0	1	4.08	2.78	0
15	NumCatalogPurchases	0	1	2.66	2.92	0
16	NumStorePurchases	0	1	5.79	3.25	0
17	NumWebVisitsMonth	0	1	5.32	2.43	0
18	AcceptedCmp3	0	1	0.0728	0.260	0
19	AcceptedCmp4	0	1	0.0746	0.263	0
20	AcceptedCmp5	0	1	0.0728	0.260	0
21	AcceptedCmp1	0	1	0.0643	0.245	0
22	AcceptedCmp2	0	1	0.0134	0.115	0
23	Complain	0	1	0.00938	0.0964	0
24	Z_CostContact	0	1	3	0	3
25	Z_Revenue	0	1	11	0	11
26	Response	0	1	0.149	0.356	0

	p25	p50	p75	p100	hist
1	2828.	5458.	8428.	11191	
2	1959	1970	1977	1996	
3	35303	51382.	68522	666666	
4	0	0	1	2	
5	0	0	1	2	
6	24	49	74	99	
7	23.8	174.	504.	1493	
8	1	8	33	199	
9	16	67	232	1725	
10	3	12	50	259	
11	1	8	33	263	

12	9	24	56	362	
13	1	2	3	15	
14	2	4	6	27	
15	0	2	4	28	
16	3	5	8	13	
17	3	6	7	20	
18	0	0	0	1	
19	0	0	0	1	
20	0	0	0	1	
21	0	0	0	1	
22	0	0	0	1	
23	0	0	0	1	
24	3	3	3	3	
25	11	11	11	11	
26	0	0	0	1	

```
Error in is.null(text_repr) || nchar(text_repr) == 0L: 'length = 17' in coercion to 'logical(1)'
Traceback:
```

(1) Data Description

0. Data Preprocessing

I downloaded the zipped data from [Kaggle](#). I extracted and saved the `.csv` file under the path `STAT301groupproject/marketing_campaign.csv`. I created a GitHub repo to store all of this work so its easier when I start working with my group-mates.

To understand the dataset structure, I first used `glimpse()` and `skim()` to inspect the total number of observations, the variable types, and basic distributional summaries (e.g., min, max, mean). This allowed me to identify the missing values in the `Income` variable and to begin categorizing variables based on their role in the analysis.

1. Dataset Summary

The **Customer Personality Analysis** dataset is a marketing dataset that contains 2,240 observations and 29 variables. Each row represents a customer, and the columns capture a wide range of information, including demographics, spending habits, campaign responses, and website interactions.

This dataset is useful for understanding customer behavior and segmenting customers for targeted marketing. For example, instead of marketing a new product to the entire customer base, a company can identify which segment is most likely to purchase and focus marketing efforts accordingly. This dataset is provided by Dr. Omar Romero-Hernandez.

Below is a grouped variable dictionary that organizes all 29 attributes into meaningful categories based on their content and analytical purpose.

Key Variables (Grouped by Category)

- 1. Customer's Information
- 2. Products (Spending in Last 2 Years)
- 3. Promotion
- 4. Place (Purchase Channels)
- 5. Other (Dummy Columns)

Full Variable Description Table (Grouped by Category)

Customer's Information

Variable	Type	Description
ID	numeric	Unique customer ID
Year_Birth	numeric	Year of birth
Education	categorical	Level of education (e.g., Graduation, PhD)
Marital_Status	categorical	Marital status
Income	numeric	Household yearly income
Kidhome	numeric	Number of children at home
Teenhome	numeric	Number of teenagers at home
Dt_Customer	datetime	Date customer enrolled
Recency	numeric	Days since last purchase
Complain	boolean	Complained in the last 2 years (1 = yes)

Products (Amount Spent in Last 2 Years)

Variable	Type	Description
MntWines	numeric	Amount spent on wine
MntFruits	numeric	Amount spent on fruits
MntMeatProducts	numeric	Amount spent on meat products
MntFishProducts	numeric	Amount spent on fish products
MntSweetProducts	numeric	Amount spent on sweet products
MntGoldProds	numeric	Amount spent on gold products

Promotion

Variable	Type	Description
NumDealsPurchases	numeric	Number of purchases using discounts
AcceptedCmp1	boolean	Accepted 1st campaign (1 = yes, 0 = no)
AcceptedCmp2	boolean	Accepted 2nd campaign (1 = yes, 0 = no)
AcceptedCmp3	boolean	Accepted 3rd campaign (1 = yes, 0 = no)
AcceptedCmp4	boolean	Accepted 4th campaign (1 = yes, 0 = no)

Variable	Type	Description
AcceptedCmp5	boolean	Accepted 5th campaign (1 = yes, 0 = no)
Response	boolean	Accepted the most recent campaign (1 = yes, 0 = no)

Place (Purchase Channels)

Variable	Type	Description
NumWebPurchases	numeric	Number of website purchases
NumCatalogPurchases	numeric	Number of catalog purchases
NumStorePurchases	numeric	Number of in-store purchases
NumWebVisitsMonth	numeric	Website visits in the last month

Other

Variable	Type	Description
Z_CostContact	numeric	Dummy cost variable (constant = 3)
Z_Revenue	numeric	Dummy revenue variable (constant = 11)

2. Missing Values & Data tidying

- The `Income` variable has **24 missing values**.
- All other variables are complete with **no missing data**.

Since the proportion of missing values here is very small, I have decided to **remove the rows with missing `Income` values**.

```
In [127... df <- df |> drop_na()
```

Note that the variables `Z_CostContact` and `Z_Revenue` have the same value for all rows, so they don't really help and I am going to drop them.

```
In [128... df <- df %>%
  select(-Z_CostContact, -Z_Revenue)
```

I am going to make **changes to some variables** so they make more sense for our analysis, in particular I will focus on

Variable(s)	Planned Transformation / Analysis	-----	-----
-----	-----		
<code>Year_Birth</code>	Convert to age	<code>Education</code>	Convert to binary: Post Graduate vs. Under Graduate
<code>Marital_Status</code>	Group into: Married vs. Not Married	<code>Kidhome</code> and <code>Teenhome</code>	Combine into a single variable: total number of children in the household
<code>Total_Spending</code>	A column that shows Total Spending = MntWines + MntFruits + MntMeatProducts + MntFishProducts + MntSweetProducts + MntGoldProds		

In [129...

```
df <- df |>

# Convert Year_Birth to Age
mutate(Age = 2025 - Year_Birth) |>

# Make Education binary : Post Graduate vs Under Graduate
mutate(Education = recode(Education,
                          'PhD' = 'Post Graduate',
                          'Master' = 'Post Graduate',
                          'Graduation' = 'Under Graduate',
                          '2n Cycle' = 'Under Graduate',
                          'Basic' = 'Under Graduate')) |>

# Make Marital_Status binary: Married vs Not Married
mutate(Marital_Status = case_when(
  Marital_Status %in% c("Married", "Together") ~ "Married", TRUE ~

# Add up Kidhome and teenhome into Num_Children
mutate(Num_Children = Kidhome + Teenhome) |>

# Add a column 'Total_Spending' by summing all spending related columns
mutate(Total_Spending = MntWines + MntFruits +MntMeatProducts + MntFi

head(df)
```

ID	Year_Birth	Education	Marital_Status	Income	Kidhome	Teenhome	Dt_Customer
<dbl>	<dbl>	<chr>	<chr>	<dbl>	<dbl>	<dbl>	<
5524	1957	Under Graduate	Not Married	58138	0	0	04-09-
2174	1954	Under Graduate	Not Married	46344	1	1	08-03-
4141	1965	Under Graduate	Married	71613	0	0	21-08-
6182	1984	Under Graduate	Married	26646	1	0	10-02-
5324	1981	Post Graduate	Married	58293	1	0	19-01-
7446	1967	Post Graduate	Married	62513	0	1	09-09-

3. Summary Statistics

I created some summary statistics tables below to help me better understand the data. Obviously there's many things we can do but I chose :

1. Summary of unique data
2. Summary of numeric columns
3. Summary of non-numeric columns

1) Summary of unique data

```
In [130... unique_df <- df |>
  summarise(across(everything(), ~ n_distinct(.))) |>
  pivot_longer(cols = everything(), names_to = "variable", values_to = "n")
  arrange(desc(n_unique))

unique_df
```


A tibble: 30 × 2

variable	n_unique
<chr>	<int>
ID	2216
Income	1974
Total_Spending	1047
MntWines	776
Dt_Customer	662
MntMeatProducts	554
MntGoldProds	212
MntFishProducts	182
MntSweetProducts	176
MntFruits	158
Recency	100
Year_Birth	59
Age	59
NumWebVisitsMonth	16
NumDealsPurchases	15
NumWebPurchases	15
NumCatalogPurchases	14
NumStorePurchases	14
Num_Children	4
Kidhome	3
Teenhome	3
Education	2
Marital_Status	2
AcceptedCmp3	2
AcceptedCmp4	2
AcceptedCmp5	2
AcceptedCmp1	2
AcceptedCmp2	2
Complain	2
Response	2

Interestingly, we have 2216 unique customers, and we see that a lot of our variables are binary.

2) Summary of numeric columns

```
In [133... numeric_df <- df |>
  select(where(is.numeric)) |>
  pivot_longer(cols = everything(), names_to = "variable", values_to = "v
  group_by(variable) |>
  summarise(
    min      = round(min(value, na.rm = TRUE), 2),
    q1       = round(quantile(value, 0.25, na.rm = TRUE), 2),
    median   = round(median(value, na.rm = TRUE), 2),
    mean     = round(mean(value, na.rm = TRUE), 2),
    q3       = round(quantile(value, 0.75, na.rm = TRUE), 2),
    max      = round(max(value, na.rm = TRUE), 2),
    sd       = round(sd(value, na.rm = TRUE), 2),
    missing  = sum(is.na(value))
  )
numeric_df
```

A tibble: 27 × 9

variable	min	q1	median	mean	q3	max	std
<chr>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>	<dbl>
AcceptedCmp1	0	0.00	0.0	0.06	0.00	1	0.24
AcceptedCmp2	0	0.00	0.0	0.01	0.00	1	0.11
AcceptedCmp3	0	0.00	0.0	0.07	0.00	1	0.25
AcceptedCmp4	0	0.00	0.0	0.07	0.00	1	0.25
AcceptedCmp5	0	0.00	0.0	0.07	0.00	1	0.25
Age	29	48.00	55.0	56.18	66.00	132	11.95
Complain	0	0.00	0.0	0.01	0.00	1	0.11
ID	0	2814.75	5458.5	5588.35	8421.75	11191	3249.35
Income	1730	35303.00	51381.5	52247.25	68522.00	666666	25173.00
Kidhome	0	0.00	0.0	0.44	1.00	2	0.95
MntFishProducts	0	3.00	12.0	37.64	50.00	259	54.77
MntFruits	0	2.00	8.0	26.36	33.00	199	39.71
MntGoldProds	0	9.00	24.5	43.97	56.00	321	51.86
MntMeatProducts	0	16.00	68.0	167.00	232.25	1725	224.22
MntSweetProducts	0	1.00	8.0	27.03	33.00	262	41.06
MntWines	0	24.00	174.5	305.09	505.00	1493	337.53
NumCatalogPurchases	0	0.00	2.0	2.67	4.00	28	2.95
NumDealsPurchases	0	1.00	2.0	2.32	3.00	15	1.96
NumStorePurchases	0	3.00	5.0	5.80	8.00	13	3.29
NumWebPurchases	0	2.00	4.0	4.09	6.00	27	2.71
NumWebVisitsMonth	0	3.00	6.0	5.32	7.00	20	2.44
Num_Children	0	0.00	1.0	0.95	1.00	3	0.71
Recency	0	24.00	49.0	49.01	74.00	99	28.90
Response	0	0.00	0.0	0.15	0.00	1	0.37
Teenhome	0	0.00	0.0	0.51	1.00	2	0.95
Total_Spending	5	69.00	396.5	607.08	1048.00	2525	602.91
Year_Birth	1893	1959.00	1970.0	1968.82	1977.00	1996	11.95

Findings:

- Income has a mean of 52,247 with max being 666,666 so we have to watch for outliers
- Wine has the highest mean spending at 305.09 followed by meat, gold, sweets, and fruits

- In store purchases are the highest mean at 5.8 followed by web purchases.

3) Summary of non-numeric columns

```
In [134]: non_numeric_df <- df |>
  select(where(~!is.numeric(.)), -Dt_Customer) |> # I took out Dt_Customer
  pivot_longer(cols = everything(), names_to = "variable", values_to = "value") |>
  group_by(variable, value) |>
  summarise(
    count = n(),
    proportion = round(count / nrow(df), 4),
    .groups = "drop"
  ) |>
  arrange(variable, desc(count))
non_numeric_df
```

A tibble: 4 × 4

	variable	value	count	proportion
	<chr>	<chr>	<int>	<dbl>
	Education	Under Graduate	1370	0.6182
	Education	Post Graduate	846	0.3818
	Marital_Status	Married	1430	0.6453
	Marital_Status	Not Married	786	0.3547

Findings : Most customers are married and relatively a lot have post graduate degrees.

(2) Question

Response Variable: `Response` is a binary variable indicating whether the customer accepted the last campaign, (1 = yes, 0 = no).

Explanatory Variables :

- `Age` — Derived from `Year_Birth`, represents the customer's age.
- `Income` — Household yearly income.
- `Education` — Converted to a binary indicator (`Postgraduate` vs. `Undergraduate`).
- `Marital_Status` — Grouped into `Married` vs. `Not Married` .
- `Complain` — Whether the customer complained in the last 2 years.
- `NumWebPurchases` — Number of purchases made through the company's website.
- `NumCatalogPurchases` — Number of purchases made through catalogs.

- `NumStorePurchases` — Number of purchases made in physical stores.
- `NumWebVisitsMonth` — Number of visits to the website in the last month.
- `Total_Spending` — Sum of spending on wine, meat, fish, fruits, sweets, and gold products over the last two years.

Hence, I can ask **How does a customer's demographic profile and spending behaviour affect likelihood of accepting a marketing campaign ?**

(3) Exploratory Data Analysis and Visualization

Lets look at an overview distribution of our data and decide what to focus on. We can use ggpairs for all numeric variables.

```
In [136... df_numeric <- df %>%
  select(Response, Age, Income, Total_Spending,
         NumWebPurchases, NumCatalogPurchases,
         NumStorePurchases, NumWebVisitsMonth)

df_numeric$Response <- as.factor(df_numeric$Response)

options(repr.plot.width = 14, repr.plot.height = 14) # setting size

ggpairs(
  data = df_numeric,
  columns = 2:8,
  aes(color = Response, alpha = 0.6), # GPT helped me with making this pr
  upper = list(continuous = wrap("cor", size = 3)),
  lower = list(continuous = wrap("points", size = 1, alpha = 0.5)),
  diag = list(continuous = wrap("densityDiag")))
```



We see the difference in distribution between people who responded to the latest marketing campaign and people who did not. `Total_Spending` is highly correlated with all three types of purchases—web, catalog, and store—suggesting multicollinearity, so we should look out for that later in the model. Additionally, customers who responded (in turquoise) tend to cluster in higher purchase and spending ranges, so there is some pattern in spending behaviour and response.

Plots we should look at :

1. Demographic vs Response

- Education
- Marital Status
- Complain

2. Numerical Predictors vs Response

- Age
- Income
- Total Spending

3. Behavioral Features vs Response

- Website Purchases
- Catalog Purchases
- Store Purchases
- Website visits


4. Correlation Between Variables

```
In [138.. # 1a. Education vs Response
edu_plot <- ggplot(df, aes(x = Education, fill = factor(Response))) +
  geom_bar(position = "fill") +
  scale_y_continuous(labels = scales::percent) +
  labs(title = "Response Rate by Education Level", y = "Proportion", fi

# 1b. Marital Status vs Response
marital_plot <- ggplot(df, aes(x = Marital_Status, fill = factor(Response)
  geom_bar(position = "fill") +
  scale_y_continuous(labels = scales::percent) +
  labs(title = "Response Rate by Marital status", y = "Proportion", fill

# 1c. Complain vs Response
complain_plot <- ggplot(df, aes(x = factor(Complain), fill = factor(Respo
  geom_bar(position = "fill") +
  scale_y_continuous(labels = scales::percent) +
  labs(title = "Respomse Rate by Complaint statuss", x = "Complain", y =

# Combined plot !
plot_grid(
  edu_plot, marital_plot, complain_plot,
  ncol = 3)
```



Customers who are married/postgraduates are a bit more likely to respond to the campaign than their counterparts. Interestingly, those who **did not complain** are marginally more responsive than those who did, so I guess satisfaction might influence engagement.


```
In [141]: # I used ChatGPT for the different colors here !
# 2a. Age vs Response
ageplot <- ggplot(df, aes(x = factor(Response), y = Age)) +
  geom_boxplot(fill = "#0073C2FF") +
  labs(title = "Age Dist. by Response", x = "Response", y = "Age")

# 2b. Income vs Response
income_plot <- ggplot(df, aes(x = factor(Response), y = Income))+
  geom_boxplot(fill = "#EFC000FF") +
  labs(title = "Income Dist. by Response",x = "Response", y = "Income")


# 2c. Total Spending vs Response
totalspend_plot <- ggplot(df, aes(x = factor(Response), y = Total_Spendin
  geom_boxplot(fill = "#868686FF") +
  labs(title = "Total Spending by Response", x = "Response", y = "Total S

# Combined plot !
plot_grid(
  ageplot,income_plot,totalspend_plot,
  ncol = 3)
```

Age Dist. by Response



Income Dist. by Response



Customers who responded to the marketing campaign tend to have higher total spending and slightly higher income than non-responders. Age distributions are the same across both groups, which means age is less important in predicting campaign response.

In [140...

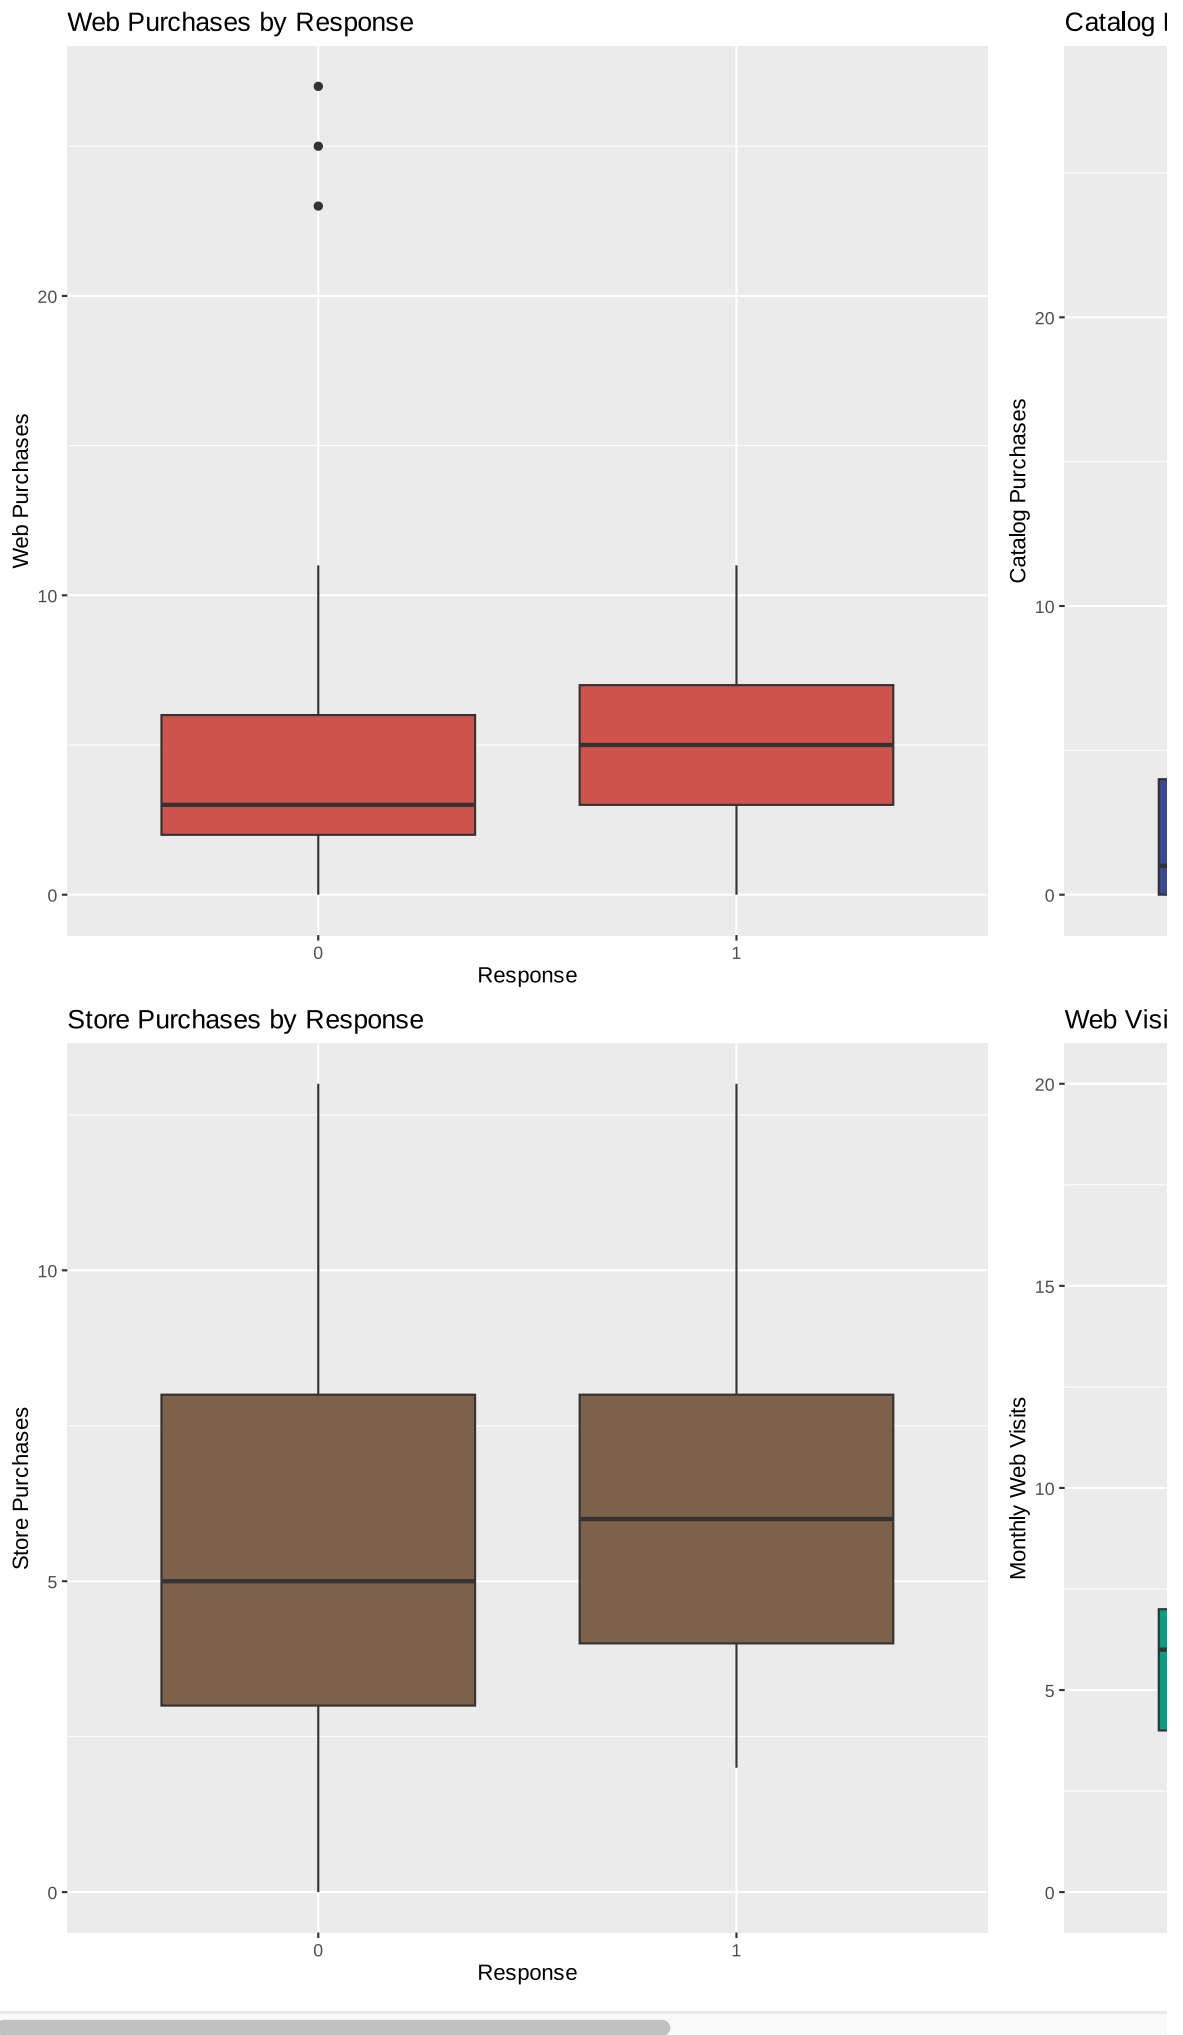
```
# I used ChatGPT for the different colors here !
# 3A. NumWebPurchases vs Response
NumWebPurchases_plot<-ggplot(df, aes(x = factor(Response), y = NumWebPurchases)) +
  geom_boxplot(fill = "#CD534CFF") +
  labs(title = "Web Purchases by Response", x = "Response", y = "Web Purchases")

# 3B. NumCatalogPurchases vs Response
NumCatalogPurchases_plot<-ggplot(df, aes(x = factor(Response), y = NumCatalogPurchases)) +
  geom_boxplot(fill = "#3B4992FF") +
  labs(title = "Catalog Purchase by Response", x = "Response", y = "Catalog Purchases")

# 3C. NumStorePurchases vs Response
NumStorePurchases_plot<-ggplot(df, aes(x = factor(Response), y = NumStorePurchases)) +
  geom_boxplot(fill = "#7E6148FF") +
  labs(title = "Store Purchases by Response", x = "Response", y = "Store Purchases")

# 3D. NumWebVisitsMonth vs Response
NumWebVisitsMonth_plot<-ggplot(df, aes(x = factor(Response), y = NumWebVisitsMonth)) +
  geom_boxplot(fill = "#00A087FF") +
  labs(title = "Web Visits by Response", x = "Response", y = "Monthly Web Visits")

NumWebPurchases_plot+NumCatalogPurchases_plot+NumStorePurchases_plot+NumWebVisitsMonth_plot
```



Customers who responded to the campaign tend to have slightly higher web and catalog purchases, suggesting a link between online/catalog buying behavior and campaign acceptance. In contrast, store purchases and web visits show little to no difference, indicating they are less important in predicting response.


4. Correlation Heatmap (predictors for customer response)

```
In [145... cor_vars <- df |>
  select(Response, Age, Income, NumWebPurchases, NumCatalogPurchases,
          NumStorePurchases, NumWebVisitsMonth, Total_Spending, Complain)

cor_matrix <- cor(cor_vars, use = "complete.obs")

ggcorrplot(cor_matrix,
  type = "lower",
  lab = TRUE,
  lab_size = 3,
  colors = c("red", "white", "blue"),
  title = "Correlation Heatmap",
  show.legend = TRUE)
```

Correlation Heatmap



From this we see that total_spending is the best predictor of response, number of catalog and web purchases also help, but there is overlap with total spending.

(4) Methods and Plan

I will use a logistic regression model to predict whether a customer will respond to a marketing campaign (Response is binary). I will use `glm()` with the `family = binomial` argument.

$$\log\left(\frac{P(\text{Response}_i = 1)}{1 - P(\text{Response}_i = 1)}\right) = \beta_0 + \sum_{j=1}^p \beta_j X_{j,i} + \sum_{\substack{t=1 \\ j \neq k}}^q \gamma_t (X_{j,i} \cdot X_{k,i})$$

Where:

- Response_i is the binary outcome (1 = responded, 0 = did not respond)
- $X_{j,i}$ are the predictor variables for observation i (listed above)
- β_0 is the intercept
- β_j are coefficients for main effects
- γ_t are coefficients for interaction effects

After fitting the full model, I will choose the variables that I find most significant and simplify the model. I will make sure my model is valid. These are the **diagnostic checks** I plan to perform:

- **Multicollinearity:** Carry out a Variance Inflation Factor (VIF). Variables with high VIFs (>5 / >10) are considered for removal to reduce redundancy and improve model interpretability.
- **Linearity of Log-Odds:** Visually verify, to ensure the relationship between continuous predictors and the logit of the response is approximately linear.
- **Independence:** Assumed based on the design of the dataset (no repeated measures or clustering).
- **Goodness-of-Fit:** Checked using a residual Q-Q plot to assess the distribution of deviance residuals. Deviations from the diagonal line would indicate model misfit.

This will improve my models predictive power.

Justification for Method

Logistic regression is suitable for a binary response variable and allows for interpretability of coefficients.

It accommodates both categorical and continuous predictors.

Adding interaction terms enables detection of heterogeneity across groups (e.g, whether marital status changes the effect of income).

Assumptions

- Observations are independent.
 - Response variable is conditionally independent of all other variables given the predictor is included in the model.
 - Log-odds of the outcome are linearly related to predictors.
 - Predictors are not strongly collinear (variables highly correlated with each other).
 - Sufficient sample size to ensure stable estimates.
-

Limitations

- Causality, this model may not be causal. What if education or marketing affects income and response ? or What if customer loyalty affects response ? But we don't record it in the dataset. So we should not make causal claims like "Higher Income causes a higher probability of response"
 - Assumes linearity in the logit, which may not hold for all variables.
 - Sensitive to outliers and influential points.
 - If the relationships we see are highly non-linear we may need another model (random forest which we haven't covered in class might be better suited)
 - Class imbalance (if present) could affect predictive performance.
-

References

imakash3011. (n.d.). Customer Personality Analysis [Data set]. Kaggle. Retrieved July 27, 2025, from <https://www.kaggle.com/datasets/imakash3011/customer-personality-analysis/data>